

# Mercari Price Suggestion Challenge

## Overview

Predicting the price of a product is a tough challenge since very similar products having minute differences such as different brand name, additional specifications, quality of the product, demand of the product, etc. can have very different prizes. For example the price of a pair of running shoes by a very common brand (say Puma) might be around INR 2,500 whereas a similar pair made by Asics can cost around INR 10,000.

Price prediction gets even more difficult when there is a huge range of products, which is common with most of the online shopping platforms. While it might be simpler to predict the price of a particular category of products using some simple criteria, it's highly challenging to predict the price of almost anything that is listed on online platforms. We may have multiple listings of the same product by a large number of sellers priced differently.

[Mercari](#) is Japan's biggest community-powered shopping app. Mercari's challenge is to build an algorithm that automatically suggests the right product prices.

## Data collection

The data can be downloaded from [Kaggle](#) competition page.

We have been provided user-inputted text descriptions of their products, including details like product category name, brand name, and item condition

## ML Problem

Using the given data, we have to come up with a model that predicts the price of a product listed on Mercari as accurately as possible.

This is a standard regression problem.

## Performance Metric

The performance of the model is measured by Root Mean Squared Logarithmic Error(RMSLE). Lesser the RMSLE, better is our prediction model.

The RMSLE is calculated as 
$$\epsilon = \sqrt{\frac{1}{n} \sum_{i=1}^n (\log(p_i + 1) - \log(a_i + 1))^2}$$

- $\epsilon$  is the RMSLE value (score)
- $n$  is the total number of observations in the (public/private) data set,
- $p_i$  is your prediction of price, and
- $a_i$  is the actual sale price for  $i$
- $\log(x)$  is the natural logarithm of  $x$

In [1]:

```
import warnings
warnings.filterwarnings('ignore')

import os
import shutil
import datetime
import gc
from tqdm import tqdm

import pandas as pd
import numpy as np
from numpy import median

%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
sns.set(style='whitegrid')
```

```

from sklearn.manifold import TSNE
from sklearn import preprocessing

from collections import Counter

import string
import re
from nltk.corpus import stopwords

import scipy
from scipy import hstack

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.metrics import mean_squared_error as mse
from math import sqrt
from sklearn.linear_model import Ridge

from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import log_loss

from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint as sp_randint
from scipy.stats import uniform
from sklearn.feature_selection.univariate_selection import SelectKBest, f_regression

```

In [2]:

```

import tracemalloc
import time

```

In [3]:

```

tracemalloc.start()

start_time = time.time()
snapshot1 = tracemalloc.take_snapshot()

```

## 0. Loading data

In [4]:

```

# https://www.kaggle.com/peterhurford/lgb-and-fm-18th-place-0-40604
def split_cat(text):
    try:
        return text.split("/")
    except:
        return ("No Label", "No Label", "No Label")

```

In [5]:

```

train = pd.read_csv('train.tsv', sep='\t',
                    dtype={'item_condition_id': 'category', 'shipping': 'category'},
                    converters={'category_name': split_cat})
test = pd.read_csv('test.tsv', sep='\t',
                   dtype={'item_condition_id': 'category', 'shipping': 'category'},
                   converters={'category_name': split_cat})

```

In [6]:

```

print('Shape of train data: ', train.shape)
print('Shape of test data: ', test.shape)
train.head(5)

```

```

Shape of train data: (1482535, 8)
Shape of test data: (3460725, 7)

```

Out [6]:

	train_id	name	item_condition_id	category_name	brand_name	price	shipping	item_description
0	0	MLB Cincinnati Reds T Shirt Size XL	3	[Men, Tops, T-shirts]	NaN	10.0	1	No description yet
1	1	Razer BlackWidow Chroma Keyboard	3	[Electronics, Computers & Tablets, Components ...]	Razer	52.0	0	This keyboard is in great condition and works ...
2	2	AVA-VIV Blouse	1	[Women, Tops & Blouses, Blouse]	Target	10.0	1	Adorable top with a hint of lace and a key hol...
3	3	Leather Horse Statues	1	[Home, Home Décor, Home Décor Accents]	NaN	35.0	1	New with tags. Leather horses. Retail for [rm]...
4	4	24K GOLD plated rose	1	[Women, Jewelry, Necklaces]	NaN	44.0	0	Complete with certificate of authenticity

## 1. Data Overview

### Handling missing values

In [7]:

```
train.isnull().any()
```

Out [7]:

```
train_id      False
name          False
item_condition_id  False
category_name  False
brand_name     True
price         False
shipping       False
item_description  True
dtype: bool
```

In [8]:

```
test.isnull().any()
```

Out [8]:

```
test_id      False
name          False
item_condition_id  False
category_name  False
brand_name     True
shipping       False
item_description  True
dtype: bool
```

Columns brand\_name, item\_description have NAs.

NAs in category\_name have been replaced by empty lists because of the converter we have used while loading the data.

We will replace NAs, empty lists with 'missing'.

In [9]:

```
# Split category_name by '/' into subcategories and replace nulls with 'missing'
train['gencat_name'] = train['category_name'].str.get(0).replace('', 'missing').astype('category')
train['subcat1_name'] = train['category_name'].str.get(1).fillna('missing').astype('category')
train['subcat2_name'] = train['category_name'].str.get(2).fillna('missing').astype('category')
train.drop('category_name', axis=1, inplace=True)
```

In [10]:

```
# Split category_name by '/' into subcategories and replace nulls with 'missing'
test['gencat_name'] = test['category_name'].str.get(0).replace('', 'missing').astype('category')
test['subcat1_name'] = test['category_name'].str.get(1).fillna('missing').astype('category')
test['subcat2_name'] = test['category_name'].str.get(2).fillna('missing').astype('category')
test.drop('category_name', axis=1, inplace=True)
```

In [11]:

```
train['item_description'].fillna('missing', inplace=True)
train['brand_name'] = train['brand_name'].fillna('missing').astype('category')
```

In [12]:

```
test['item_description'].fillna('missing', inplace=True)
test['brand_name'] = test['brand_name'].fillna('missing').astype('category')
```

## Check for duplicate rows, NAs

In [13]:

```
train[train.duplicated()]
```

Out[13]:

train_id	name	item_condition_id	brand_name	price	shipping	item_description	gencat_name	subcat1_name	subcat2_name
----------	------	-------------------	------------	-------	----------	------------------	-------------	--------------	--------------

No duplicate rows in train data

In [14]:

```
train.isnull().any()
```

Out[14]:

```
train_id      False
name          False
item_condition_id  False
brand_name    False
price         False
shipping      False
item_description False
gencat_name   False
subcat1_name  False
subcat2_name  False
dtype: bool
```

## Check for rows with invalid price

In [15]:

```
print('Removed {} rows'.format(len(train[train.price<=0])))
train = train[train.price > 0].reset_index(drop=True)
```

Removed 874 rows

## Column-wise overview of data

name

In [16]:

```
train.name.describe()
```

Out[16]:

```
count      1481661
unique      1224596
top         Bundle
freq         2232
Name: name, dtype: object
```

### item\_condition\_id

In [17]:

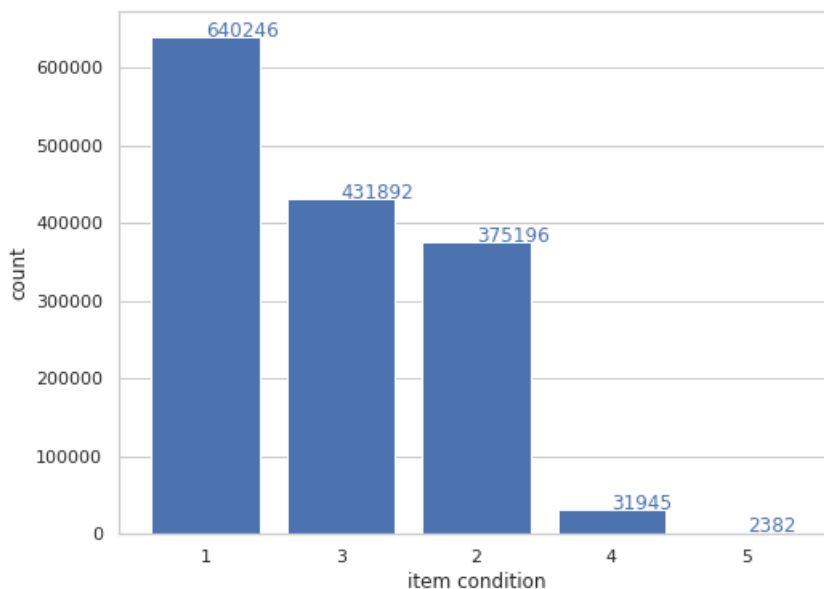
```
train.item_condition_id.describe()
```

Out[17]:

```
count      1481661
unique         5
top          1
freq      640246
Name: item_condition_id, dtype: object
```

In [18]:

```
condition_count = Counter(list(train.item_condition_id))
x, y = zip(*condition_count.most_common())
plt.figure(figsize=[8,6])
plt.bar(x, y, )
for i, val in enumerate(y):
    plt.annotate(val, (x[i], y[i]), color='b')
plt.xlabel('item condition')
plt.ylabel('count')
plt.grid(False, axis='x')
plt.show()
```



Majority of the items are in **condition 1**

### brand\_name

In [19]:

```
train.brand_name.describe()
```

Out[19]:

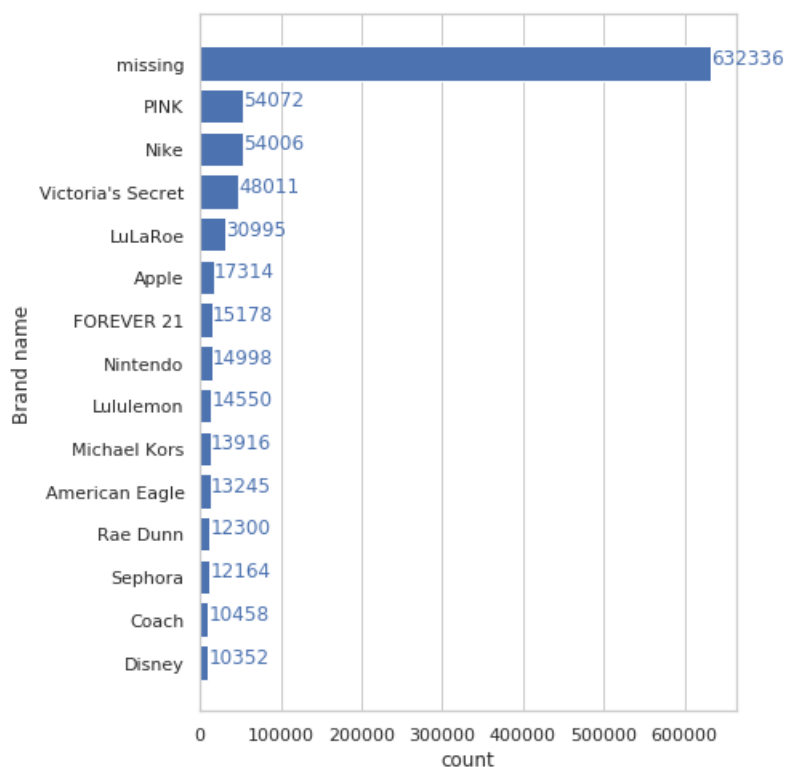
Out[19]:

```
count      1481661
unique       4808
top         missing
freq        632336
Name: brand_name, dtype: object
```

In [20]:

```
brand_count = Counter(list(train.brand_name.values))
x, y = zip(*brand_count.most_common(15))

plt.figure(figsize=[6,8])
plt.barh(x, y)
for i, val in enumerate(y):
    plt.annotate(val, (y[i], x[i]), color='b')
plt.gca().invert_yaxis()
plt.ylabel('Brand name')
plt.xlabel('count')
plt.grid(False, axis='y')
plt.show()
```



In [21]:

```
brand_missing = train[train.brand_name=='missing'].shape[0]
print('Brand name is missing for {} datapoints, i.e. {:.2f} % of train data.'
      .format(brand_missing, 100.0*brand_missing/train.shape[0]))
```

Brand name is missing for 632336 datapoints, i.e. 42.68 % of train data.

## gencat\_name

In [22]:

```
train.gencat_name.describe()
```

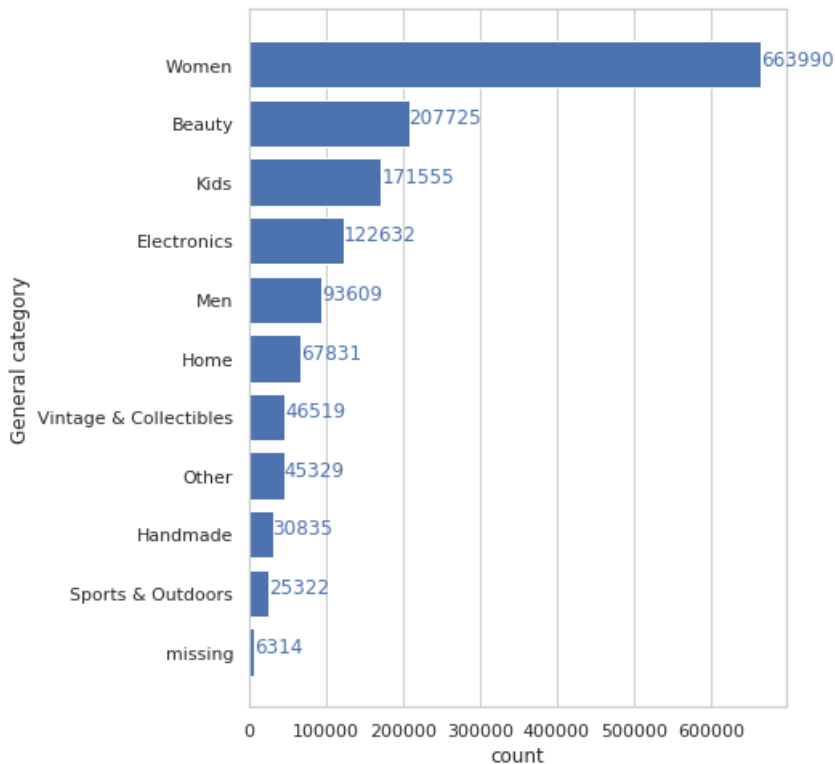
Out[22]:

```
count      1481661
unique        11
top         Women
freq        663990
```

```
Name: gencat_name, dtype: object
```

```
In [23]:
```

```
gencat_count = Counter(list(train.gencat_name.values))
x, y = zip(*gencat_count.most_common(15))
plt.figure(figsize=[6,8])
plt.barh(x, y)
for i, val in enumerate(y):
    plt.annotate(val, (y[i], x[i]), color='b')
plt.gca().invert_yaxis()
plt.ylabel('General category')
plt.xlabel('count')
plt.grid(False, axis='y')
plt.show()
```



Majority of the items are from the category **women**

```
In [24]:
```

```
gencat_missing = train[train.gencat_name=='missing'].shape[0]
print('category name is missing for {} datapoints, i.e. {:.2f} % of train data.'
      .format(gencat_missing, 100.0*gencat_missing/train.shape[0]))
```

category name is missing for 6314 datapoints, i.e. 0.43 % of train data.

**subcat1\_name**

```
In [25]:
```

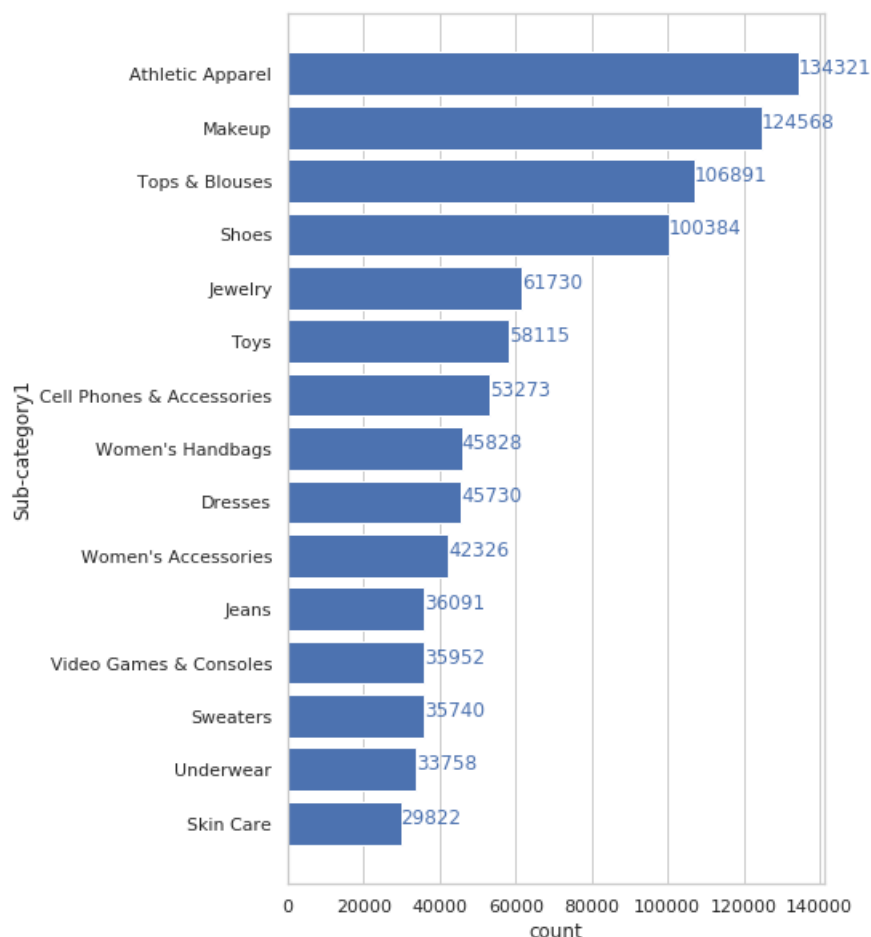
```
train.subcat1_name.describe()
```

```
Out[25]:
```

```
count          1481661
unique           114
top      Athletic Apparel
freq           134321
Name: subcat1_name, dtype: object
```

In [26]:

```
subcat1_count = Counter(list(train.subcat1_name.values))
x, y = zip(*subcat1_count.most_common(15))
plt.figure(figsize=[6,10])
plt.barh(x, y)
for i, val in enumerate(y):
    plt.annotate(val, (y[i], x[i]), color='b')
plt.gca().invert_yaxis()
plt.ylabel('Sub-category1')
plt.xlabel('count')
plt.grid(False, axis='y')
plt.show()
```



In [27]:

```
subcat1_missing = train[train.subcat1_name=='missing'].shape[0]
print('subcategory1 name is missing for {} datapoints, i.e. {:.2f} % of train data.'
      .format(subcat1_missing, 100.0*subcat1_missing/train.shape[0]))
```

subcategory1 name is missing for 6314 datapoints, i.e. 0.43 % of train data.

## subcat2\_name

In [28]:

```
train.subcat2_name.describe()
```

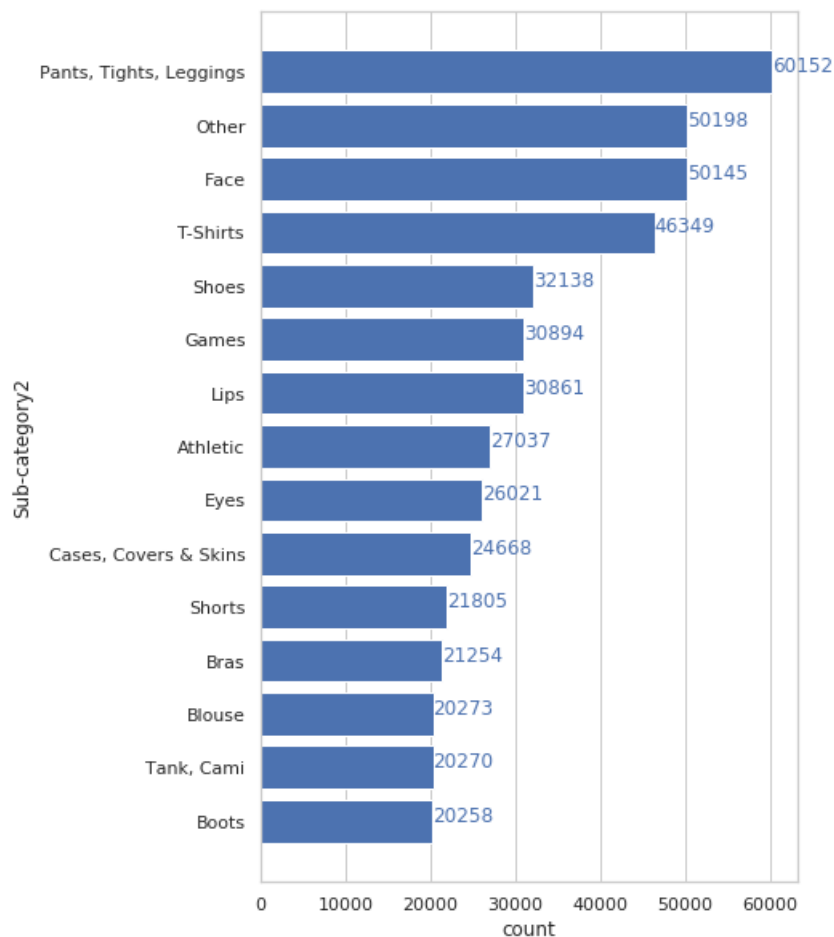
Out[28]:

```
count          1481661
unique           871
top      Pants, Tights, Leggings
freq           60152
Name: subcat2_name, dtype: object
```



In [29]:

```
subcat2_count = Counter(list(train.subcat2_name.values))
x, y = zip(*subcat2_count.most_common(15))
plt.figure(figsize=[6,10])
plt.barh(x, y)
for i, val in enumerate(y):
    plt.annotate(val, (y[i], x[i]), color='b')
plt.gca().invert_yaxis()
plt.ylabel('Sub-category2')
plt.xlabel('count')
plt.grid(False, axis='y')
plt.show()
```



In [30]:

```
subcat2_missing = train[train.subcat2_name=='missing'].shape[0]
print('subcategory2 name is missing for {} datapoints, i.e. {:.2f} % of train data.'
      .format(subcat2_missing, 100.0*subcat2_missing/train.shape[0]))
```

subcategory2 name is missing for 6314 datapoints, i.e. 0.43 % of train data.

## item\_description

In [31]:

```
desc_missing = train[train.item_description=='missing'].shape[0]
print('item description is missing for {} datapoints, i.e. {:.5f} % of train data.'
      .format(desc_missing, 100.0*desc_missing/train.shape[0]))
```

item description is missing for 4 datapoints, i.e. 0.00027 % of train data.

In [32]:

```
train[train.item_description=='missing']
```

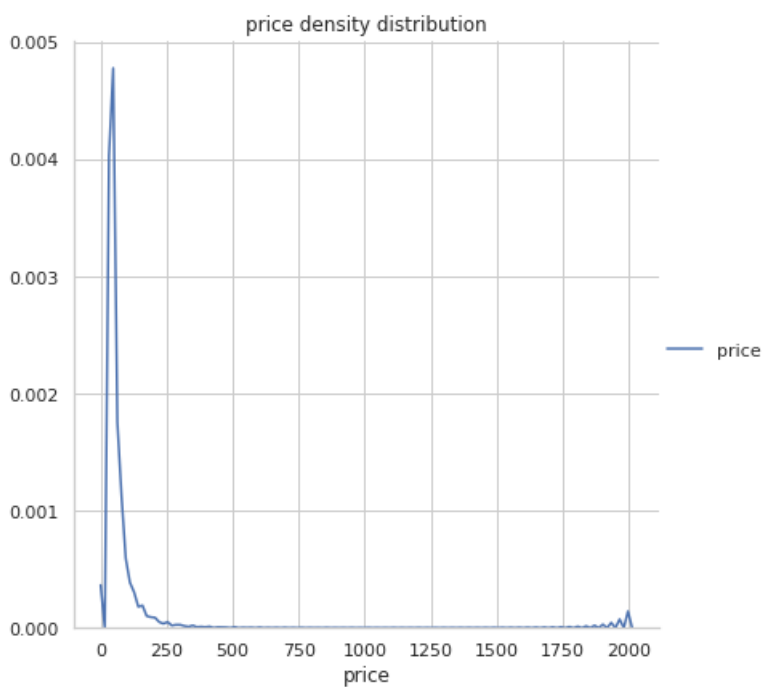
Out[32]:

	train_id	name	item_condition_id	brand_name	price	shipping	item_description	gencat_name	subcat1_name	subcat2_name
511264	511535	Shoes for Michelle	4	missing	9.0	0	missing	Kids	Girls 0-24 Mos	St
860756	861230	Lipgloss	4	missing	49.0	0	missing	Beauty	Makeup	
1224217	1224924	Disney Minnie Head band	3	Disney	9.0	0	missing	Women	Women's Accessories	Accesso
1263521	1264242	For Bianca	3	missing	10.0	1	missing	Women	Women's Accessories	Scarve W

price

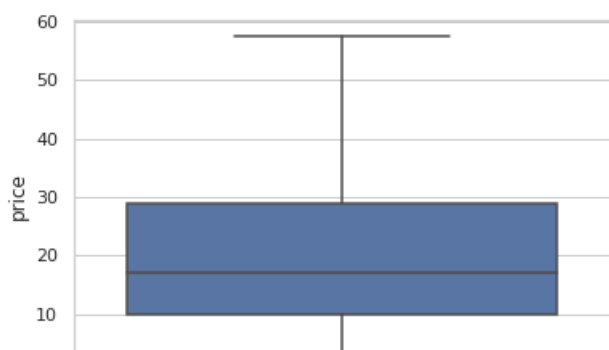
In [33]:

```
sns.FacetGrid(train,size=6) \
    .map(sns.kdeplot,"price") \
    .add_legend();
plt.title('price density distribution')
plt.show();
```



In [34]:

```
sns.boxplot(y='price', data=train, showfliers=False)
plt.show()
```



In [35]:

```
for i in range(0, 100, 10):
    var = train["price"].values
    var = np.sort(var, axis = None)
    print("{} percentile value is {}".format(i, var[int(len(var)*(float(i)/100))]))
print("100 percentile value is ", var[-1])
```

```
0 percentile value is 3.0
10 percentile value is 7.0
20 percentile value is 10.0
30 percentile value is 12.0
40 percentile value is 14.0
50 percentile value is 17.0
60 percentile value is 20.0
70 percentile value is 26.0
80 percentile value is 34.0
90 percentile value is 51.0
100 percentile value is 2009.0
```

In [36]:

```
for i in range(90, 100, 1):
    var = train["price"].values
    var = np.sort(var, axis = None)
    print("{} percentile value is {}".format(i, var[int(len(var)*(float(i)/100))]))
print("100 percentile value is ", var[-1])
```

```
90 percentile value is 51.0
91 percentile value is 55.0
92 percentile value is 58.0
93 percentile value is 62.0
94 percentile value is 67.0
95 percentile value is 75.0
96 percentile value is 85.0
97 percentile value is 99.0
98 percentile value is 122.0
99 percentile value is 170.0
100 percentile value is 2009.0
```

**97% of datapoints have price less than USD 100. Very few (only 1%) datapoints have price more than USD 170**

## 2. Pre-processing

In [37]:

```
def preprocess_name(text_col):
    preprocessed_names = []
    for sentence in tqdm(text_col.values):
        sent = sentence.replace('\\r', ' ')
        sent = sent.replace('\\\"', ' ')
        sent = sent.replace('\\n', ' ')
        sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
        preprocessed_names.append(sent.lower().strip())
    return preprocessed_names

stopwords = stopwords.words('english')
def preprocess_desc(text_col):
    preprocessed_descs = []
    for sentence in tqdm(text_col.values):
        sent = sentence.replace('\\r', ' ')
        sent = sent.replace('\\\"', ' ')
        sent = sent.replace('\\n', ' ')
        sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
        sent = ' '.join(e for e in sent.split() if e not in stopwords)
        preprocessed_descs.append(sent.lower().strip())
    return preprocessed_descs
```

In [38]:

```
train['preprocessed_name'] = preprocess_name(train['name'])
test['preprocessed_name'] = preprocess_name(test['name'])

train['preprocessed_description'] = preprocess_desc(train['item_description'])
test['preprocessed_description'] = preprocess_desc(test['item_description'])
```

```
100%|██████████| 1481661/1481661 [00:11<00:00, 124510.07it/s]
100%|██████████| 3460725/3460725 [00:27<00:00, 124025.17it/s]
100%|██████████| 1481661/1481661 [02:33<00:00, 9672.48it/s]
100%|██████████| 3460725/3460725 [05:58<00:00, 9656.55it/s]
```

In [39]:

```
def clean_cat(cat_values):
    '''takes categorical column values as arguments and returns list of cleaned categories'''

    categories = list(cat_values)

    cat_list = []
    for i in tqdm(categories):
        i = re.sub('[^A-Za-z0-9]+', ' ', i)
        i = i.replace(' ', '')
        i = i.replace('&', '_')
        cat_list.append(i.strip())

    return cat_list
```

In [40]:

```
train['gencat_name'] = clean_cat(train['gencat_name'].values)
test['gencat_name'] = clean_cat(test['gencat_name'].values)

train['subcat1_name'] = clean_cat(train['subcat1_name'].values)
test['subcat1_name'] = clean_cat(test['subcat1_name'].values)

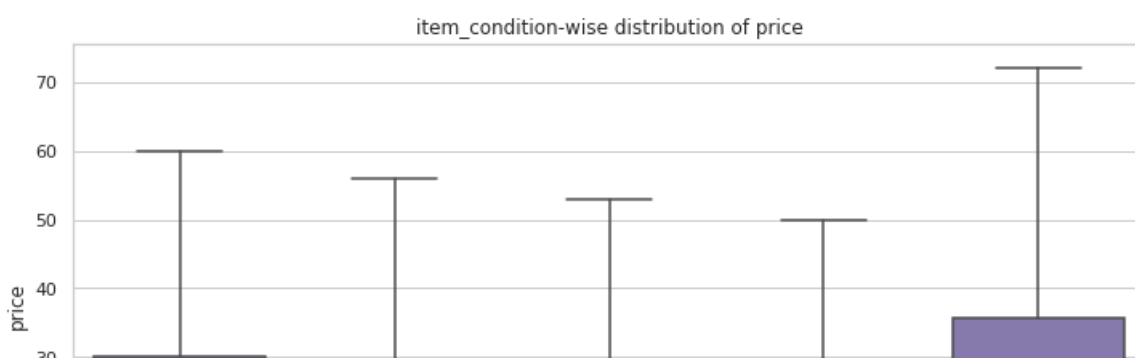
train['subcat2_name'] = clean_cat(train['subcat2_name'].values)
test['subcat2_name'] = clean_cat(test['subcat2_name'].values)
```

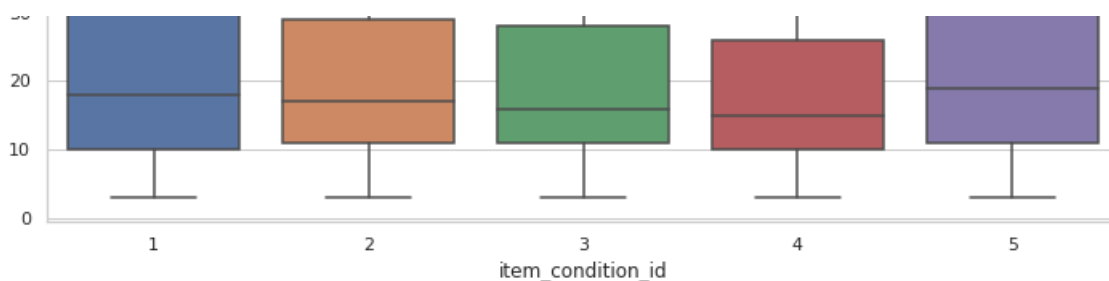
```
100%|██████████| 1481661/1481661 [00:05<00:00, 295231.76it/s]
100%|██████████| 3460725/3460725 [00:11<00:00, 289588.24it/s]
100%|██████████| 1481661/1481661 [00:06<00:00, 219928.05it/s]
100%|██████████| 3460725/3460725 [00:15<00:00, 219610.83it/s]
100%|██████████| 1481661/1481661 [00:06<00:00, 222446.74it/s]
100%|██████████| 3460725/3460725 [00:15<00:00, 221882.62it/s]
```

### 3. Basic Exploratory Data Analysis

In [38]:

```
sns.set(style='whitegrid')
plt.figure(figsize=(12,6))
sns.boxplot(x='item_condition_id', y='price', data=train, showfliers=False)
plt.title('item_condition-wise distribution of price')
plt.show()
```

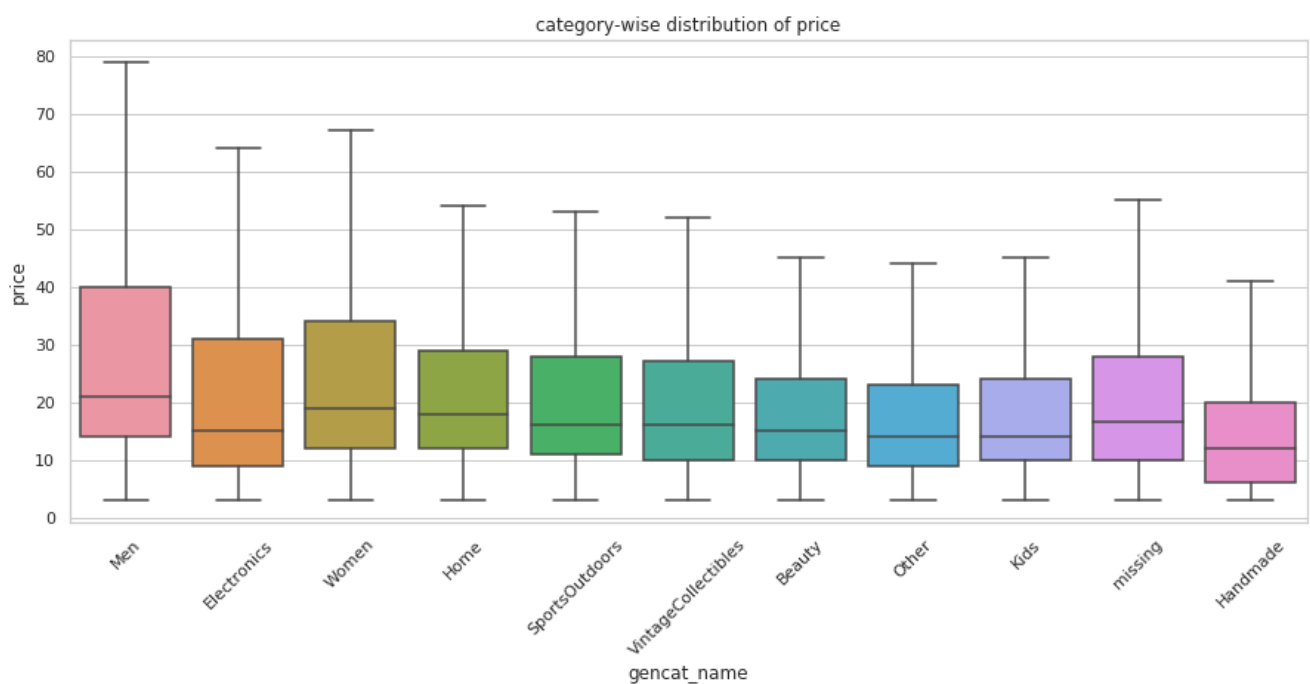




There is slight variation of price based on item condition. Median Price decreases as we go from condition 1 to 4. Items in condition 5 seem to be having higher price, which is a bit strange.

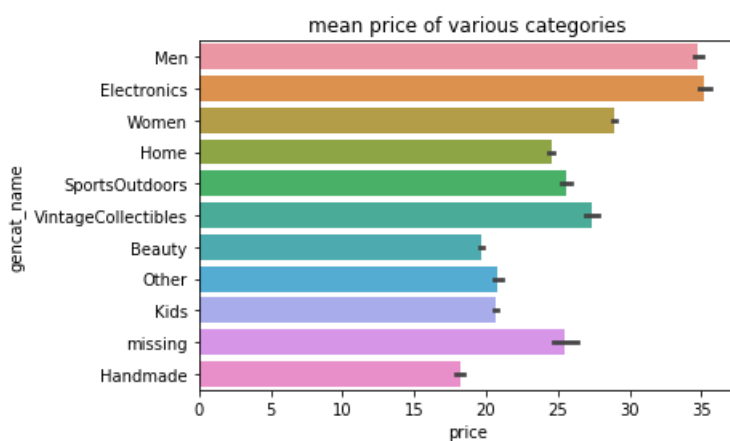
In [39]:

```
plt.figure(figsize=(15,6))
sns.boxplot(y='price', x='gencat_name', data=train, showfliers=False)
plt.xticks(rotation=45)
plt.title('category-wise distribution of price')
plt.show()
```



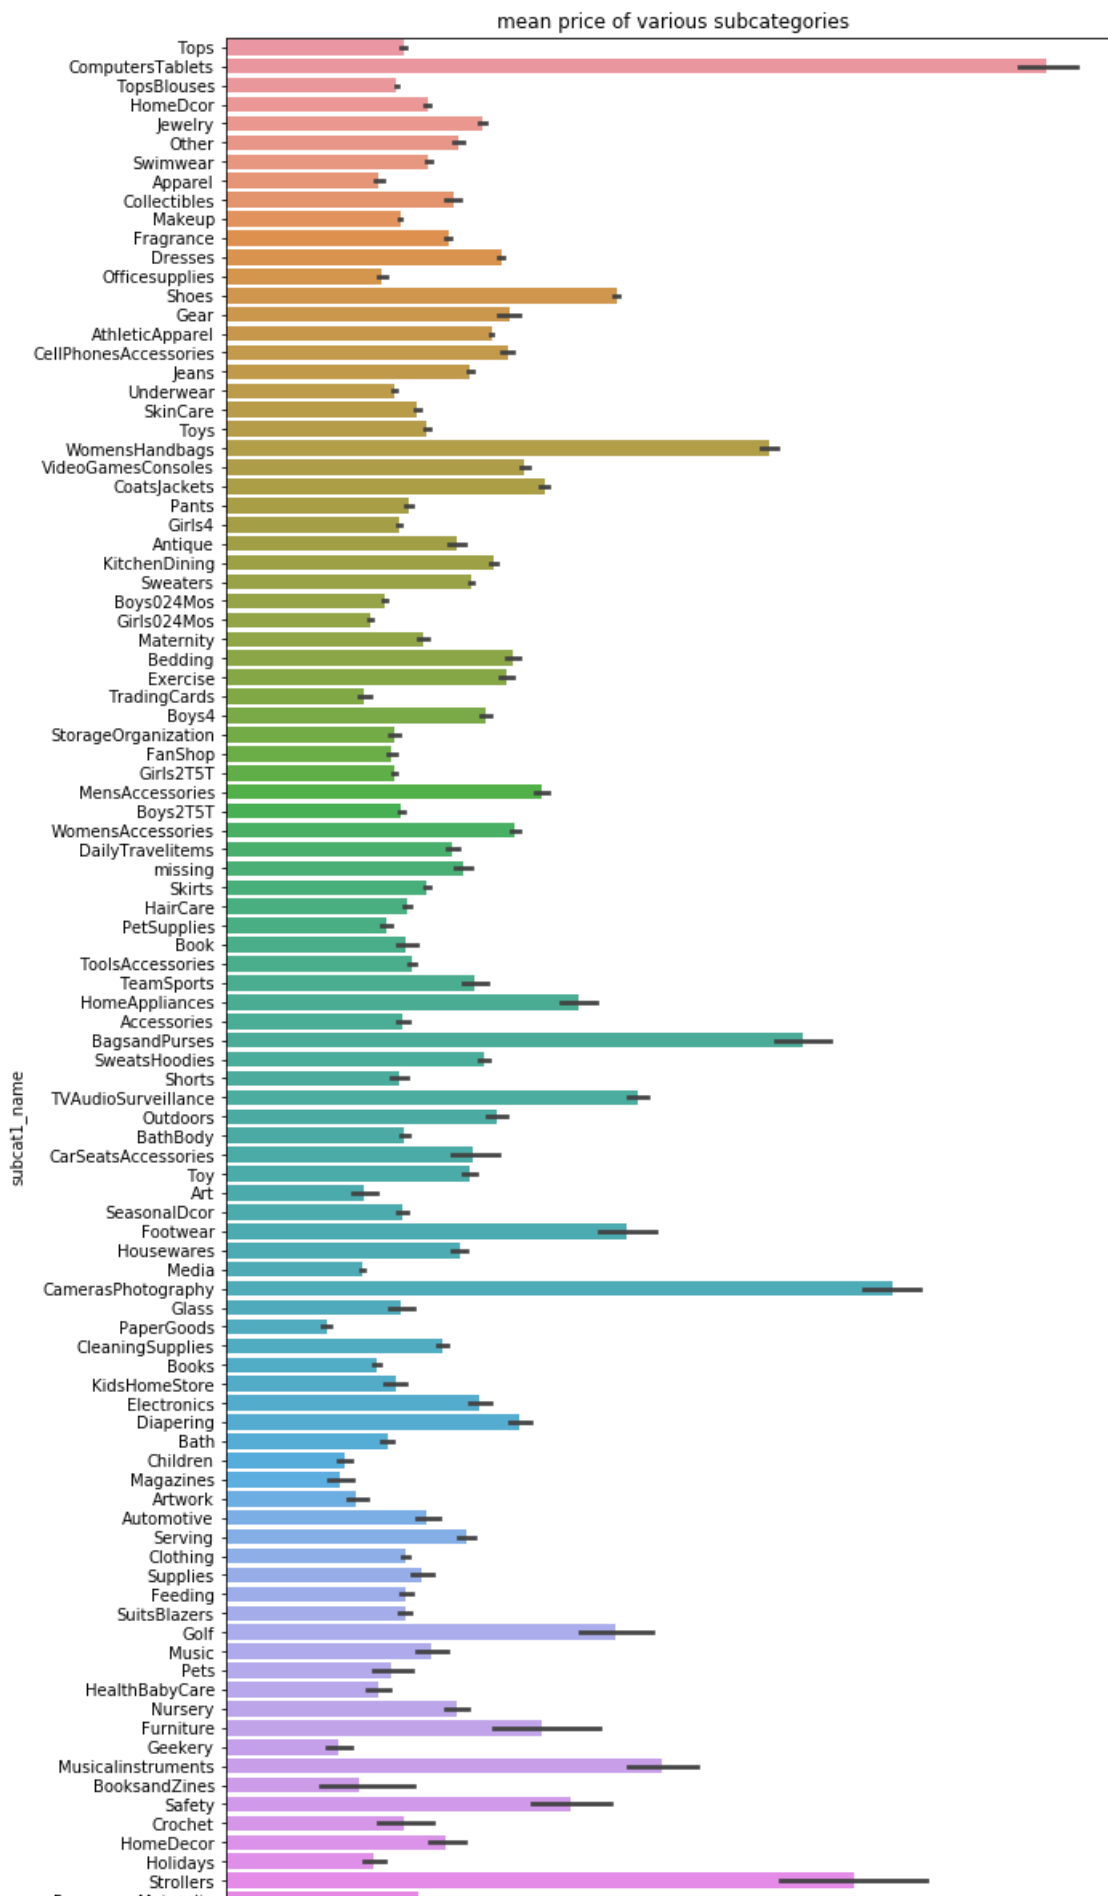
In [38]:

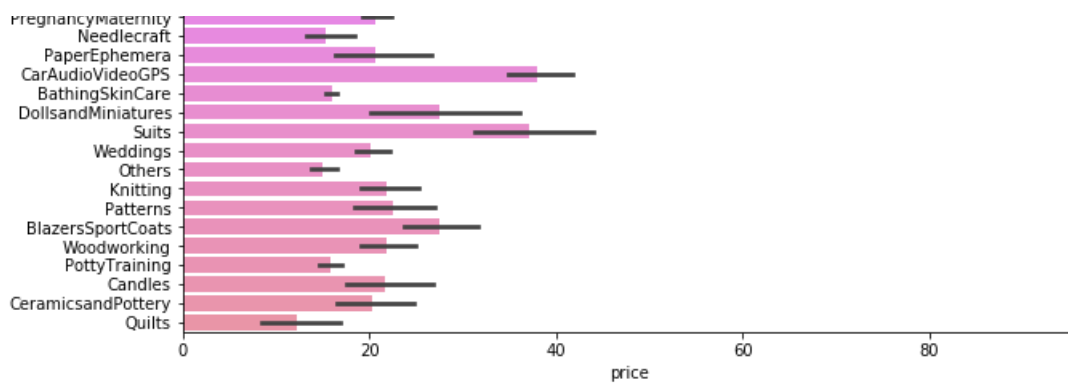
```
sns.barplot(y='gencat_name', x='price', data=train)
plt.title('mean price of various categories')
plt.show()
```



In [39]:

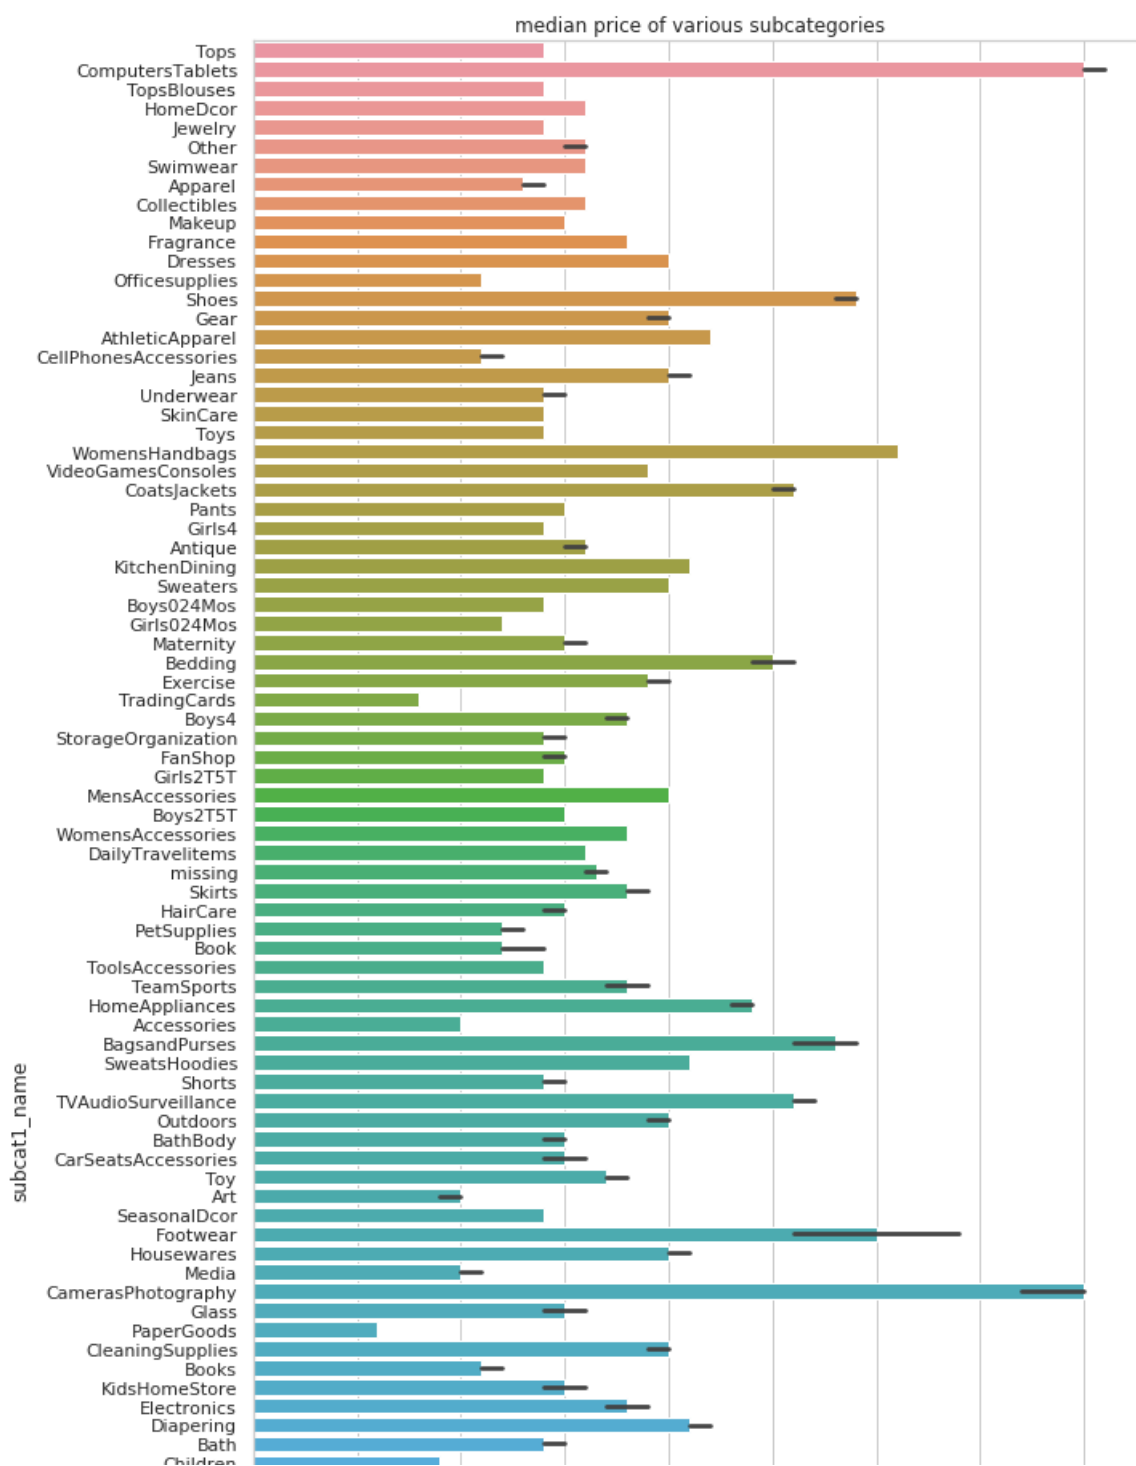
```
plt.figure(figsize=(10,25))
sns.barplot(y='subcat1_name', x='price', data=train)
plt.title('mean price of various subcategories')
plt.show()
```

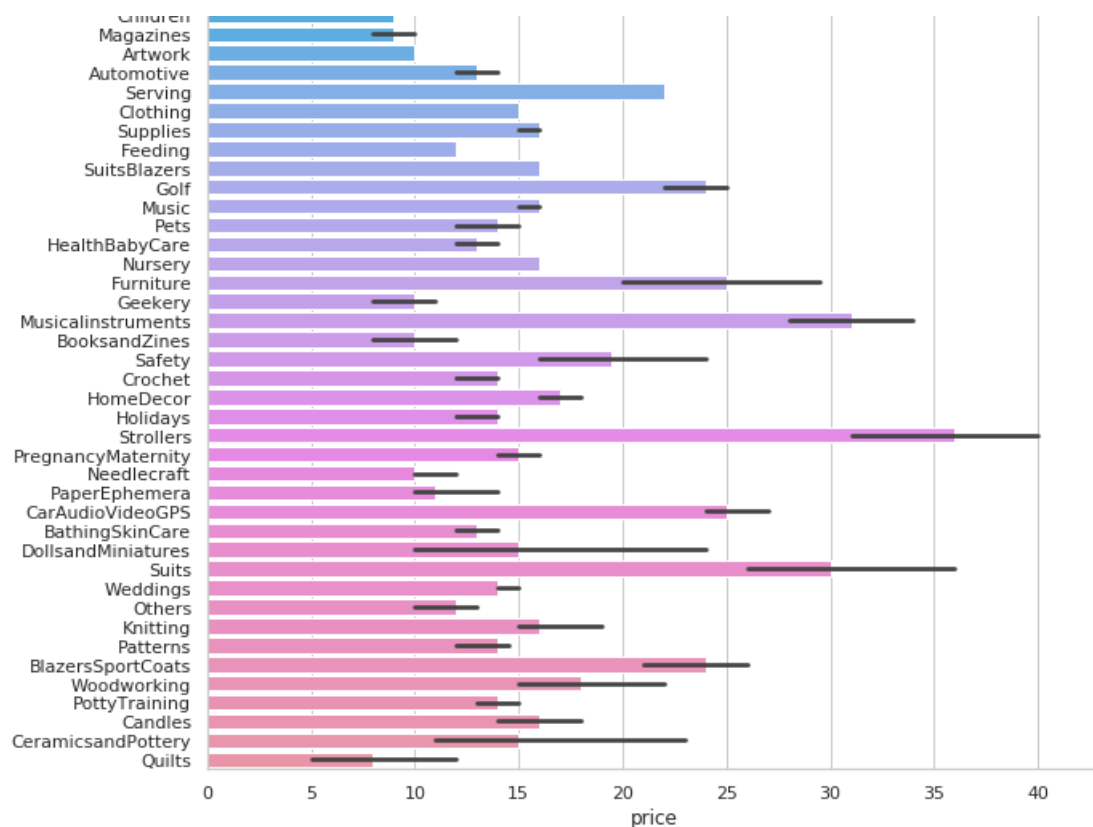




In [41]:

```
plt.figure(figsize=(10,25))
sns.barplot(y='subcat1_name', x='price', data=train, estimator=median)
plt.title('median price of various subcategories')
plt.show()
```





Prices of items belonging to various categories and subcategories vary significantly. This indicates that categories are going to be important features in determining the price of an item.

## 4. Featurization

### 4.1. Feature Engineering

**name\_first:** cleaned name first word

**name\_first\_count:** count of name in data

**gencat\_name\_count:** count of gencat in data

**subcat1\_name\_count:** count of subcat1 in data

**subcat2\_name\_count:** count of subcat2 in data

**brand\_name\_count:** count of brand\_name in data

In [41]:

```
def get_name_first(name):
    name = re.sub('[^A-Za-z0-9]+', ' ', name).split()
    if len(name):
        return name[0].lower()
    return ''

train['name_first'] = train['name'].apply(get_name_first)
test['name_first'] = test['name'].apply(get_name_first)
```

In [42]:

```
def transform_test(base_col, feat_col):
    """
    Returns feat_col column of test data by mapping from the values already calculated for the same column in train data
```



*Parameters:*

*base\_col: column based on which a transform(count, mean, median) has been applied*

*feat\_col: desired feature column after applying the transform*

*'''*

*#Create dictionary of feature values from train data*

```
di = pd.Series(train[feat_col].values, index=train[base_col].values).to_dict()
```

*#Map test data using dictionary and fill NAs with 0*

```
if base_col == 'item_condition_id':
```

```
    #No chance of NAs
```

```
    return test[base_col].map(di).astype(float)
```

```
return test[base_col].map(di).fillna(0)
```

In [43]:

```
train['name_first_count'] = train.groupby('name_first')['name_first'].transform('count')
```

```
test['name_first_count'] = transform_test('name_first', 'name_first_count')
```

```
train['gencat_name_count'] = train.groupby('gencat_name')['gencat_name'].transform('count')
```

```
test['gencat_name_count'] = transform_test('gencat_name', 'gencat_name_count')
```

```
train['subcat1_name_count'] = train.groupby('subcat1_name')['subcat1_name'].transform('count')
```

```
test['subcat1_name_count'] = transform_test('subcat1_name', 'subcat1_name_count')
```

```
train['subcat2_name_count'] = train.groupby('subcat2_name')['subcat2_name'].transform('count')
```

```
test['subcat2_name_count'] = transform_test('subcat2_name', 'subcat2_name_count')
```

```
train['brand_name_count'] = train.groupby('brand_name')['brand_name'].transform('count')
```

```
test['brand_name_count'] = transform_test('brand_name', 'brand_name_count')
```

**NameLower:** # lowercase letters in name

**DescriptionLower:** # lowercase letters in description

**NameUpper:** # uppercase letters in name

**DescriptionUpper:** # uppercase letters in description

**name\_len:** char length of name

**des\_len:** char length of desc

**name\_desc\_len\_ratio:** name\_len / des\_len

**desc\_word\_count**

**mean\_des:** 10 \* desc\_word\_count / des\_len

**name\_word\_count**

**mean\_name:** 10 \* name\_word\_count / name\_len

**desc\_letters\_per\_word:** des\_len / desc\_word\_count

**name\_letters\_per\_word:** name\_len / name\_word\_count

**NameLowerRatio:** NameLower / name\_len

**DescriptionLowerRatio:** DescriptionLower / des\_len

**NameUpperRatio:** NameUpper / name\_len

**DescriptionUpperRatio:** DescriptionUpper / des\_len

In [44]:

```
train['NameLower'] = train.name.str.count('[a-z]')
```

```
train['DescriptionLower'] = train.item_description.str.count('[a-z]')
```

```
train['NameUpper'] = train.name.str.count('[A-Z]')
```

```
train['DescriptionUpper'] = train.item_description.str.count('[A-Z]')
```

```
train['name_len'] = train['name'].apply(lambda x: len(x))
```

```

train['name_len'] = train['name'].apply(lambda x: len(x))
train['des_len'] = train['item_description'].apply(lambda x: len(x))
train['name_desc_len_ratio'] = train['name_len']/train['des_len']
train['desc_word_count'] = train['item_description'].apply(lambda x: len(x.split()))
train['mean_des'] = train['item_description'].apply(lambda x: 0 if len(x) == 0 else float(len(x.split())) / len(x)) * 10
train['name_word_count'] = train['name'].apply(lambda x: len(x.split()))
train['mean_name'] = train['name'].apply(lambda x: 0 if len(x) == 0 else float(len(x.split())) / len(x)) * 10
train['desc_letters_per_word'] = train['des_len'] / train['desc_word_count']
train['name_letters_per_word'] = train['name_len'] / train['name_word_count']
train['NameLowerRatio'] = train['NameLower'] / train['name_len']
train['DescriptionLowerRatio'] = train['DescriptionLower'] / train['des_len']
train['NameUpperRatio'] = train['NameUpper'] / train['name_len']
train['DescriptionUpperRatio'] = train['DescriptionUpper'] / train['des_len']

```

In [45]:

```

test['NameLower'] = test.name.str.count('[a-z]')
test['DescriptionLower'] = test.item_description.str.count('[a-z]')
test['NameUpper'] = test.name.str.count('[A-Z]')
test['DescriptionUpper'] = test.item_description.str.count('[A-Z]')
test['name_len'] = test['name'].apply(lambda x: len(x))
test['des_len'] = test['item_description'].apply(lambda x: len(x))
test['name_desc_len_ratio'] = test['name_len']/test['des_len']
test['desc_word_count'] = test['item_description'].apply(lambda x: len(x.split()))
test['mean_des'] = test['item_description'].apply(lambda x: 0 if len(x) == 0 else float(len(x.split())) / len(x)) * 10
test['name_word_count'] = test['name'].apply(lambda x: len(x.split()))
test['mean_name'] = test['name'].apply(lambda x: 0 if len(x) == 0 else float(len(x.split())) / len(x)) * 10
test['desc_letters_per_word'] = test['des_len'] / test['desc_word_count']
test['name_letters_per_word'] = test['name_len'] / test['name_word_count']
test['NameLowerRatio'] = test['NameLower'] / test['name_len']
test['DescriptionLowerRatio'] = test['DescriptionLower'] / test['des_len']
test['NameUpperRatio'] = test['NameUpper'] / test['name_len']
test['DescriptionUpperRatio'] = test['DescriptionUpper'] / test['des_len']

```

**NamePunctCount:** # punctuations in name

**DescriptionPunctCount:** # punctuations in desc

**NamePunctCountRatio:** NamePunctCount / name\_word\_count

**DescriptionPunctCountRatio:** DescriptionPunctCount / desc\_word\_count

**NameDigitCount:** # digits in name

**DescriptionDigitCount:** # digits in desc

**NameDigitCountRatio:** NameDigitCount / name\_word\_count

**DescriptionDigitCountRatio:** DescriptionDigitCount / desc\_word\_count

**stopword\_ratio\_desc:** # stopwords in desc / desc\_word\_count

**num\_sum:** Sum of numbers in desc

**weird\_characters\_desc:** # non-alphanumeric, non-punct in desc

**weird\_characters\_name:** # non-alphanumeric, non-punct in desc

**prices\_count:** # of [rm] (removed price) in desc

**price\_in\_name:** 1 if desc contains [rm]; 0 otherwise

In [46]:

```

from nltk.corpus import stopwords

RE_PUNCTUATION = '|'.join([re.escape(x) for x in string.punctuation])
s_words = {x: 1 for x in stopwords.words('english')} #converting to dictionary for fast look up
non_alphanumpunct = re.compile(u'^A-Za-z0-9\.\?!,\; \(\)\[\]\\"\\$'+')

```

In [47]:

```
#https://www.kaggle.com/peterhurford/lgb-and-fm-18th-place-0-40604
```

```
def to_number(x):
    try:
        if not x.isdigit():
            return 0
        x = int(x)
        if x > 100:
            return 100
        else:
            return x
    except:
        return 0
```

In [48]:

```
train['NamePunctCount'] = train.name.str.count(RE_PUNCTUATION)
train['DescriptionPunctCount'] = train.item_description.str.count(RE_PUNCTUATION)
train['NamePunctCountRatio'] = train['NamePunctCount'] / train['name_word_count']
train['DescriptionPunctCountRatio'] = train['DescriptionPunctCount'] / train['desc_word_count']
train['NameDigitCount'] = train.name.str.count('[0-9]')
train['DescriptionDigitCount'] = train.item_description.str.count('[0-9]')
train['NameDigitCountRatio'] = train['NameDigitCount'] / train['name_word_count']
train['DescriptionDigitCountRatio'] = train['DescriptionDigitCount']/train['desc_word_count']
train['stopword_ratio_desc'] = train['item_description'].apply(lambda x: len([w for w in x.split()
if w in s_words])) / train['desc_word_count']
train['num_sum'] = train['item_description'].apply(lambda x: sum([to_number(s) for s in x.split()])
)
train['weird_characters_desc'] = train['item_description'].str.count(non_alphanumpunct)
train['weird_characters_name'] = train['name'].str.count(non_alphanumpunct)
train['prices_count'] = train['item_description'].str.count('[rm]')
train['price_in_name'] = train['item_description'].str.contains('[rm]', regex=False).astype('category')
```

In [49]:

```
test['NamePunctCount'] = test.name.str.count(RE_PUNCTUATION)
test['DescriptionPunctCount'] = test.item_description.str.count(RE_PUNCTUATION)
test['NamePunctCountRatio'] = test['NamePunctCount'] / test['name_word_count']
test['DescriptionPunctCountRatio'] = test['DescriptionPunctCount'] / test['desc_word_count']
test['NameDigitCount'] = test.name.str.count('[0-9]')
test['DescriptionDigitCount'] = test.item_description.str.count('[0-9]')
test['NameDigitCountRatio'] = test['NameDigitCount'] / test['name_word_count']
test['DescriptionDigitCountRatio'] = test['DescriptionDigitCount']/test['desc_word_count']
test['stopword_ratio_desc'] = test['item_description'].apply(lambda x: len([w for w in x.split() if
w in s_words])) / test['desc_word_count']
test['num_sum'] = test['item_description'].apply(lambda x: sum([to_number(s) for s in x.split()]))
test['weird_characters_desc'] = test['item_description'].str.count(non_alphanumpunct)
test['weird_characters_name'] = test['name'].str.count(non_alphanumpunct)
test['prices_count'] = test['item_description'].str.count('[rm]')
test['price_in_name'] = test['item_description'].str.contains('[rm]', regex=False).astype('category')
```

**brand\_mean\_price:** mean price of items by a brand

**name\_mean\_price:** mean price of an item by name

**gencat\_mean\_price:** mean price of items belonging to gencat

**subcat1\_mean\_price:** mean price of items belonging to subcat1

**subcat2\_mean\_price:** mean price of items belonging to subcat2

**condition\_mean\_price:** mean price of items by condition

**brand\_median\_price:** median price of items by a brand

**name\_median\_price:** median price of an item by name

**gencat\_median\_price:** median price of items belonging to gencat

**subcat1\_median\_price:** median price of items belonging to subcat1

**subcat2\_median\_price:** median price of items belonging to subcat2

**condition\_median\_price:** median price of items by condition

In [50]:

```
train['brand_mean_price'] = train.groupby('brand_name')['price'].transform('mean')
test['brand_mean_price'] = transform_test('brand_name', 'brand_mean_price')

train['name_mean_price'] = train.groupby('name_first')['price'].transform('mean')
test['name_mean_price'] = transform_test('name_first', 'name_mean_price')

train['gencat_mean_price'] = train.groupby('gencat_name')['price'].transform('mean')
test['gencat_mean_price'] = transform_test('gencat_name', 'gencat_mean_price')

train['subcat1_mean_price'] = train.groupby('subcat1_name')['price'].transform('mean')
test['subcat1_mean_price'] = transform_test('subcat1_name', 'subcat1_mean_price')

train['subcat2_mean_price'] = train.groupby('subcat2_name')['price'].transform('mean')
test['subcat2_mean_price'] = transform_test('subcat2_name', 'subcat2_mean_price')

train['condition_mean_price'] = train.groupby('item_condition_id')['price'].transform('mean')
test['condition_mean_price'] = transform_test('item_condition_id', 'condition_mean_price')
```

In [51]:

```
train['brand_median_price'] = train.groupby('brand_name')['price'].transform('median')
test['brand_median_price'] = transform_test('brand_name', 'brand_median_price')

train['name_median_price'] = train.groupby('name_first')['price'].transform('median')
test['name_median_price'] = transform_test('name_first', 'name_median_price')

train['gencat_median_price'] = train.groupby('gencat_name')['price'].transform('median')
test['gencat_median_price'] = transform_test('gencat_name', 'gencat_median_price')

train['subcat1_median_price'] = train.groupby('subcat1_name')['price'].transform('median')
test['subcat1_median_price'] = transform_test('subcat1_name', 'subcat1_median_price')

train['subcat2_median_price'] = train.groupby('subcat2_name')['price'].transform('median')
test['subcat2_median_price'] = transform_test('subcat2_name', 'subcat2_median_price')

train['condition_median_price'] = train.groupby('item_condition_id')['price'].transform('median')
test['condition_median_price'] = transform_test('item_condition_id', 'condition_median_price')
```

In [52]:

```
train.drop(['name', 'item_description'], axis=1, inplace=True)
test.drop(['name', 'item_description'], axis=1, inplace=True)
```

In [53]:

```
print(train.shape, test.shape)
```

```
(1481661, 59) (3460725, 58)
```

## Univariate analysis on the above features

In [61]:

```
plt.figure(figsize=(18,18))

plt.subplot(3,3,1)
sns.regplot(x='brand_mean_price', y='price', data=train, scatter_kws={'alpha':0.3}, line_kws={'color':'orange'})
plt.title('brand_mean_price vs price(target)')

plt.subplot(3,3,2)
sns.regplot(x='gencat_mean_price', y='price', data=train, scatter_kws={'alpha':0.3}, line_kws={'color':'orange'})
```

```

plt.title('category_mean_price vs price(target)')

plt.subplot(3,3,3)
sns.regplot(x='subcat1_mean_price', y='price', data=train, scatter_kws={'alpha':0.3}, line_kws={'color':'orange'})
plt.title('subcategory_mean_price vs price(target)')

plt.subplot(3,3,4)
sns.regplot(x='subcat2_mean_price', y='price', data=train, scatter_kws={'alpha':0.3}, line_kws={'color':'orange'})
plt.title('subcategory_mean_price vs price(target)')

plt.subplot(3,3,5)
sns.regplot(x='condition_mean_price', y='price', data=train, scatter_kws={'alpha':0.3}, line_kws={'color':'orange'})
plt.title('condition_mean_price vs price(target)')

plt.subplot(3,3,6)
sns.regplot(x='brand_median_price', y='price', data=train, scatter_kws={'alpha':0.3}, line_kws={'color':'orange'})
plt.title('brand_median_price vs price(target)')

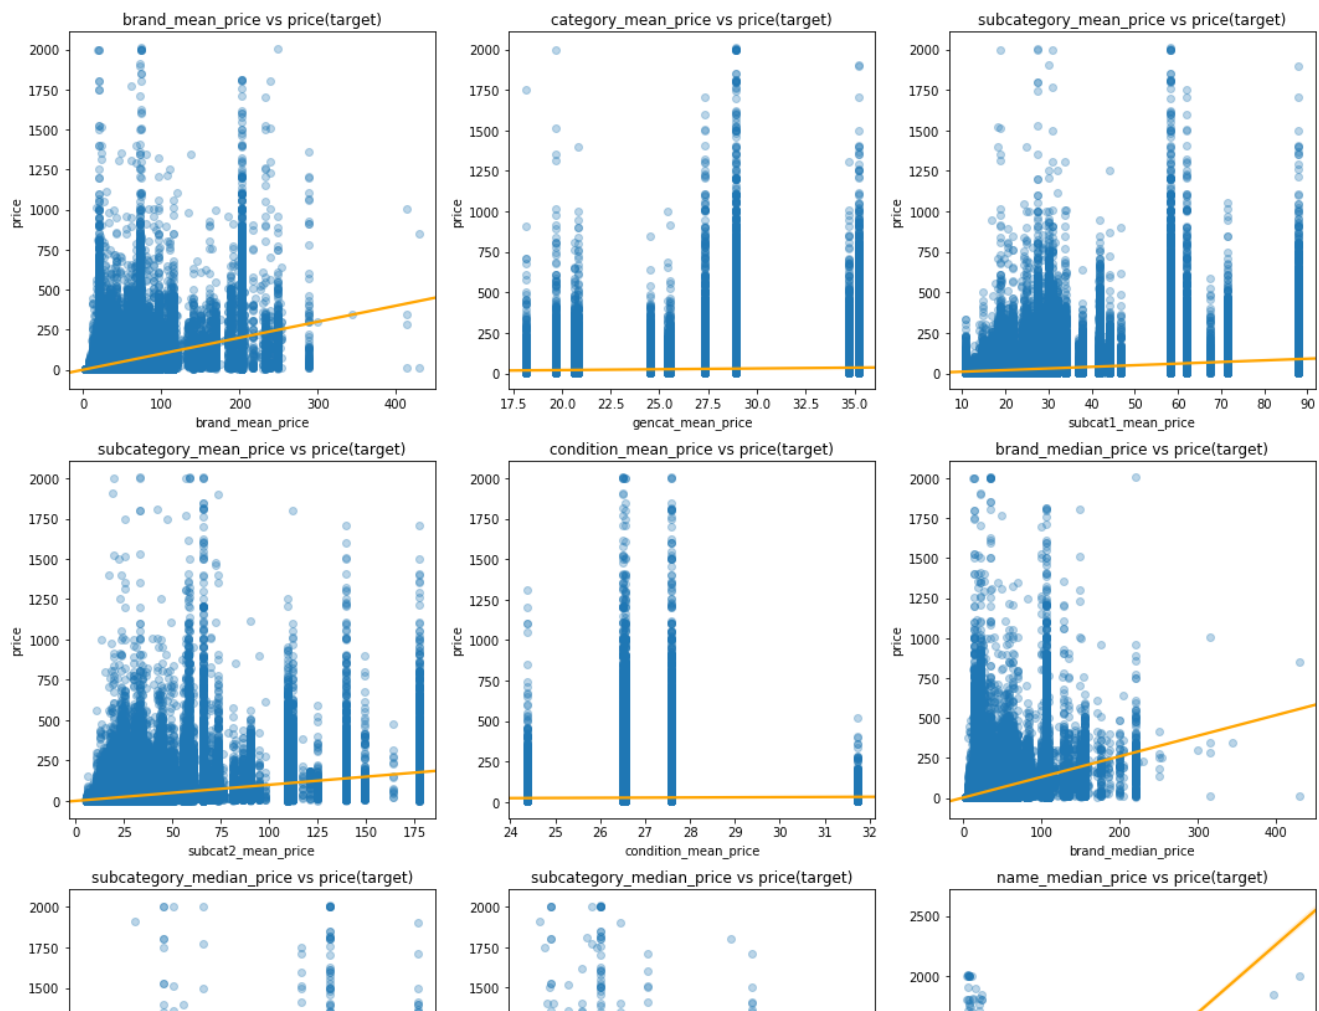
plt.subplot(3,3,7)
sns.regplot(x='subcat1_median_price', y='price', data=train, scatter_kws={'alpha':0.3}, line_kws={'color':'orange'})
plt.title('subcategory_median_price vs price(target)')

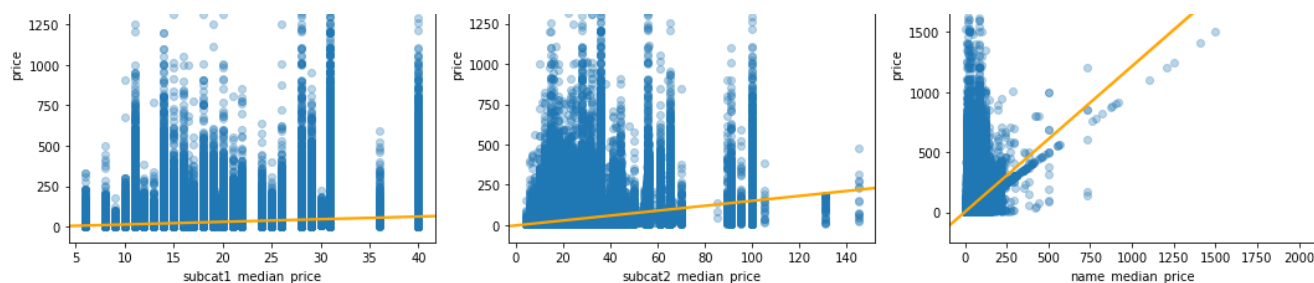
plt.subplot(3,3,8)
sns.regplot(x='subcat2_median_price', y='price', data=train, scatter_kws={'alpha':0.3}, line_kws={'color':'orange'})
plt.title('subcategory_median_price vs price(target)')

plt.subplot(3,3,9)
sns.regplot(x='name_median_price', y='price', data=train, scatter_kws={'alpha':0.3}, line_kws={'color':'orange'})
plt.title('name_median_price vs price(target)')

plt.show()

```





Features such as **brand\_mean\_price**, **brand\_median price**, **subcat2\_mean\_price**, **subcat2\_median\_price** show a linear trend therefore seem to be useful in determining price of items.

### Dropping rows with blank name and description

In [54]:

```
n_rows = train.shape[0]
train = train[train.preprocessed_name != ''].reset_index(drop=True)

print('Dropped {} rows'.format(n_rows - train.shape[0]))
```

Dropped 19 rows

In [55]:

```
n_rows = train.shape[0]
train = train[train.preprocessed_description != ''].reset_index(drop=True)

print('Dropped {} rows'.format(n_rows - train.shape[0]))

print('Shape of train data: ', train.shape)
```

Dropped 567 rows

Shape of train data: (1481075, 59)

## 4.2. Train, Test split for cross validation

In [56]:

```
from sklearn.model_selection import train_test_split

y_tr = np.log1p(train['price'])
train.drop(['price'], axis=1, inplace=True)

train_df, cv_df, y_train, y_cv = train_test_split(train, y_tr, test_size=0.1, random_state=42)

print('Train size: {}, CV size: {}, Test size: {}'.format(train_df.shape, cv_df.shape, test.shape))
```

Train size: (1332967, 58), CV size: (148108, 58), Test size: (3460725, 58)

In [57]:

```
del train, y_tr
gc.collect()
```

Out[57]:

120

## 4.3. Categorical features

One-hot encoding of brand\_name, gencat\_name, subcat1\_name, subcat2\_name.

In [58]:

```
#Cleaning brand name before using count vectorizer
# Using same preprocessing as used earlier for categories: 'clean_cat()' function

train_df['brand_name'] = clean_cat(train_df['brand_name'].values)
cv_df['brand_name'] = clean_cat(cv_df['brand_name'].values)
test['brand_name'] = clean_cat(test['brand_name'].values)
```

```
100%|██████████| 1332967/1332967 [00:05<00:00, 255964.09it/s]
100%|██████████| 148108/148108 [00:00<00:00, 251876.77it/s]
100%|██████████| 3460725/3460725 [00:13<00:00, 256942.02it/s]
```

In [59]:

```
vectorizer = CountVectorizer(lowercase=False, binary=True)
train_brand_oneHot = vectorizer.fit_transform(train_df['brand_name'].values)

cv_brand_oneHot = vectorizer.transform(cv_df['brand_name'].values)
test_brand_oneHot = vectorizer.transform(test['brand_name'].values)

print("Shape of matrices after one hot encoding")
print(train_brand_oneHot.shape, "\n", cv_brand_oneHot.shape, "\n", test_brand_oneHot.shape)
```

```
Shape of matrices after one hot encoding
(1332967, 4651)
(148108, 4651)
(3460725, 4651)
```

In [60]:

```
vectorizer = CountVectorizer(lowercase=False, binary=True)
train_gencat_oneHot = vectorizer.fit_transform(train_df['gencat_name'].values)

cv_gencat_oneHot = vectorizer.transform(cv_df['gencat_name'].values)
test_gencat_oneHot = vectorizer.transform(test['gencat_name'].values)

print("Shape of matrices after one hot encoding")
print(train_gencat_oneHot.shape, "\n", cv_gencat_oneHot.shape, "\n", test_gencat_oneHot.shape)
```

```
Shape of matrices after one hot encoding
(1332967, 11)
(148108, 11)
(3460725, 11)
```

In [61]:

```
vectorizer = CountVectorizer(lowercase=False, binary=True)
train_subcat1_oneHot = vectorizer.fit_transform(train_df['subcat1_name'].values)

cv_subcat1_oneHot = vectorizer.transform(cv_df['subcat1_name'].values)
test_subcat1_oneHot = vectorizer.transform(test['subcat1_name'].values)

print("Shape of matrices after one hot encoding")
print(train_subcat1_oneHot.shape, "\n", cv_subcat1_oneHot.shape, "\n", test_subcat1_oneHot.shape)
```

```
Shape of matrices after one hot encoding
(1332967, 114)
(148108, 114)
(3460725, 114)
```

In [62]:

```
vectorizer = CountVectorizer(lowercase=False, binary=True)
train_subcat2_oneHot = vectorizer.fit_transform(train_df['subcat2_name'].values)

cv_subcat2_oneHot = vectorizer.transform(cv_df['subcat2_name'].values)
test_subcat2_oneHot = vectorizer.transform(test['subcat2_name'].values)

print("Shape of matrices after one hot encoding")
```

```
print(train_subcat2_oneHot.shape, "\n", cv_subcat2_oneHot.shape, "\n", test_subcat2_oneHot.shape)
```

Shape of matrices after one hot encoding  
(1332967, 863)  
(148108, 863)  
(3460725, 863)

## 4.4. Tfidf vectorization on text features

1-3 grams of name

1-3 grams of item\_description

In [63]:

```
vectorizer = TfidfVectorizer(ngram_range=(1, 3), min_df=3, max_features=250000)

train_name_tfidf = vectorizer.fit_transform(train_df['preprocessed_name'].values)

cv_name_tfidf = vectorizer.transform(cv_df['preprocessed_name'].values)
test_name_tfidf = vectorizer.transform(test['preprocessed_name'].values)

print("Shape of matrices after vectorization")
print(train_name_tfidf.shape, "\n", cv_name_tfidf.shape, "\n", test_name_tfidf.shape)
```

Shape of matrices after vectorization  
(1332967, 250000)  
(148108, 250000)  
(3460725, 250000)

In [64]:

```
vectorizer = TfidfVectorizer(ngram_range=(1, 3), min_df=5, max_features=500000)

train_description_tfidf = vectorizer.fit_transform(train_df['preprocessed_description'].values)

cv_description_tfidf = vectorizer.transform(cv_df['preprocessed_description'].values)
test_description_tfidf = vectorizer.transform(test['preprocessed_description'].values)

print("Shape of matrices after vectorization")
print(train_description_tfidf.shape, "\n", cv_description_tfidf.shape, "\n", test_description_tfidf.shape)
```

Shape of matrices after vectorization  
(1332967, 500000)  
(148108, 500000)  
(3460725, 500000)

## 5. Data preparation

In [65]:

```
submission_df = pd.DataFrame(test['test_id'])
print(submission_df.shape)
submission_df.head()
```

(3460725, 1)

Out[65]:

test_id	
0	0
1	1
2	2
3	3



## 5.1. Normalize numerical features

In [66]:

```
cols = set(train_df.columns.values) - {'train_id'}
skip_cols = {'preprocessed_name', 'item_condition_id', 'brand_name',
            'shipping', 'preprocessed_description', 'gencat_name',
            'subcat1_name', 'subcat2_name', 'name_first', 'price_in_name'}

cols_to_normalize = cols - skip_cols
print("Normalizing following columns: ", cols_to_normalize)

def normalize(df):
    result1 = df.copy()
    for feature_name in df.columns:
        if (feature_name in cols_to_normalize):
            max_value = df[feature_name].max()
            min_value = df[feature_name].min()
            result1[feature_name] = (df[feature_name] - min_value) / (max_value - min_value)
    return result1
```

Normalizing following columns: {'NameLower', 'name\_len', 'name\_word\_count', 'weird\_characters\_name', 'NamePunctCount', 'name\_median\_price', 'name\_first\_count', 'des\_len', 'condition\_mean\_price', 'subcat1\_median\_price', 'desc\_word\_count', 'prices\_count', 'num\_sum', 'NameLowerRatio', 'condition\_median\_price', 'desc\_letters\_per\_word', 'DescriptionDigitCount', 'DescriptionUpperRatio', 'name\_letters\_per\_word', 'subcat2\_median\_price', 'subcat2\_name\_count', 'DescriptionLowerRatio', 'NameDigitCountRatio', 'weird\_characters\_desc', 'brand\_name\_count', 'gencat\_median\_price', 'DescriptionLower', 'DescriptionPunctCount', 'NameUpper', 'brand\_median\_price', 'brand\_mean\_price', 'gencat\_name\_count', 'subcat1\_name\_count', 'name\_desc\_len\_ratio', 'mean\_name', 'subcat1\_mean\_price', 'stopword\_ratio\_desc', 'DescriptionPunctCountRatio', 'NamePunctCountRatio', 'DescriptionDigitCountRatio', 'mean\_des', 'DescriptionUpper', 'name\_mean\_price', 'gencat\_mean\_price', 'NameDigitCount', 'NameUpperRatio', 'subcat2\_mean\_price'}

In [67]:

```
train_normalized = normalize(train_df)
cv_normalized = normalize(cv_df)
test_normalized = normalize(test)
```

In [68]:

```
del train_df, cv_df, test
gc.collect()
```

Out[68]:

20

## 5.2. Remove non-features from dataframes

In [69]:

```
#Separating and storing all numerical features

X_tr = train_normalized[list(cols_to_normalize)]
X_val = cv_normalized[list(cols_to_normalize)]
X_te = test_normalized[list(cols_to_normalize)]

X_tr.head(2)
```

Out[69]:

	NameLower	name_len	name_word_count	weird_characters_name	NamePunctCount	name_median_price	name_first_count	c
660536	0.131579	0.166667	0.0625	0.0	0.0	0.008513	0.144537	0.

2 rows × 47 columns

In [70]:

```

from scipy.sparse import csr_matrix

# Storing categorical features to sparse matrix

X_tr_cat = csr_matrix(pd.get_dummies(train_normalized[['item_condition_id', 'shipping', 'price_in_name']], sparse=True).values)

X_cv_cat = csr_matrix(pd.get_dummies(cv_normalized[['item_condition_id', 'shipping', 'price_in_name']], sparse=True).values)

X_te_cat = csr_matrix(pd.get_dummies(test_normalized[['item_condition_id', 'shipping', 'price_in_name']], sparse=True).values)

print(X_tr_cat.shape, X_cv_cat.shape, X_te_cat.shape)

```

(1332967, 9) (148108, 9) (3460725, 9)

In [71]:

```

del train_normalized, cv_normalized, test_normalized
gc.collect()

```

Out[71]:

20

### 5.3. Consolidate all features to a sparse matrix

In [72]:

```

from scipy.sparse import hstack

# stack all categorical and text sparse matrices

train_sparse = hstack((train_brand_oneHot, train_gencat_oneHot, train_subcat1_oneHot,
train_subcat2_oneHot, \
train_name_tfidf, train_description_tfidf, X_tr_cat)).tocsr()

cv_sparse = hstack((cv_brand_oneHot, cv_gencat_oneHot, cv_subcat1_oneHot, cv_subcat2_oneHot, \
cv_name_tfidf, cv_description_tfidf, X_cv_cat)).tocsr()

test_sparse = hstack((test_brand_oneHot, test_gencat_oneHot, test_subcat1_oneHot,
test_subcat2_oneHot, \
test_name_tfidf, test_description_tfidf, X_te_cat)).tocsr()

```

In [73]:

```

print(train_sparse.shape, cv_sparse.shape, test_sparse.shape)

```

(1332967, 755648) (148108, 755648) (3460725, 755648)

In [74]:

```

# stack dense feature matrix with categorical and text vectors

X_train = hstack((X_tr.values, train_sparse)).tocsr()

X_cv = hstack((X_val.values, cv_sparse)).tocsr()

X_test = hstack((X_te.values, test_sparse)).tocsr()

```

In [75]:

```
In [75]:
```

```
print('Train size: {}, CV size: {}, Test size: {}'.format(X_train.shape, X_cv.shape, X_test.shape)
)
```

Train size: (1332967, 755695), CV size: (148108, 755695), Test size: (3460725, 755695)

```
In [76]:
```

```
del vectorizer
del train_brand_oneHot, train_gencat_oneHot, train_subcat1_oneHot, train_subcat2_oneHot, \
     train_name_tfidf, train_description_tfidf, X_tr_cat

del cv_brand_oneHot, cv_gencat_oneHot, cv_subcat1_oneHot, cv_subcat2_oneHot, cv_name_tfidf,
cv_description_tfidf, X_cv_cat

del test_brand_oneHot, test_gencat_oneHot, test_subcat1_oneHot, test_subcat2_oneHot, \
     test_name_tfidf, test_description_tfidf, X_te_cat
gc.collect()
```

```
Out[76]:
```

100

## 6. Modeling

```
In [77]:
```

```
from sklearn.metrics import mean_squared_error as mse
from math import sqrt
from sklearn.linear_model import Ridge
```

Target variable is  $\ln(1+y)$ , therefore calculating mean square error on  $\ln(1+y)$  will effectively give us MSLE on  $y$

- $\text{MSLE}(y^{\wedge}, y) = \text{MSE}(\ln(1+y^{\wedge}), \ln(1+y))$

### 6.1. Ridge Model

Linear least squares with l2 regularization [sklearn.linear\\_model.Ridge](#)

#### Hyper-parameter tuning

```
In [79]:
```

```
alpha = [1, 2, 3, 3.5, 4, 4.5, 5, 6, 7]
cv_rmsle_array=[]
for i in tqdm(alpha):
    model = Ridge(solver="sag", random_state=42, alpha=i)
    model.fit(X_train, y_train)
    preds_cv = model.predict(X_cv)
    cv_rmsle_array.append(sqrt(mse(y_cv, preds_cv)))

for i in range(len(cv_rmsle_array)):
    print ('RMSLE for alpha = ', alpha[i], 'is', cv_rmsle_array[i])

best_alpha = np.argmin(cv_rmsle_array)

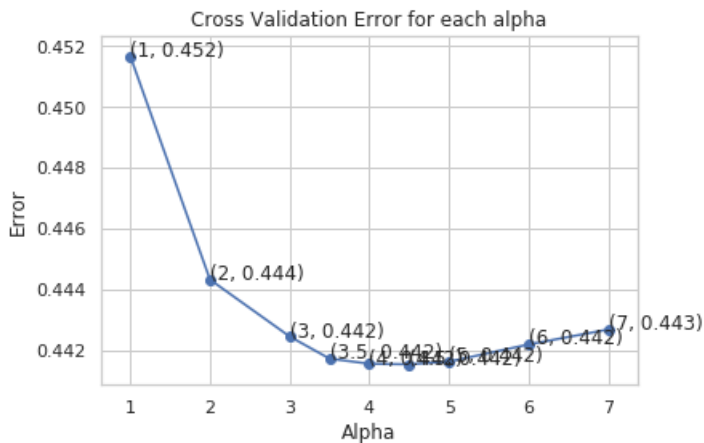
fig, ax = plt.subplots()
ax.plot(alpha, cv_rmsle_array)
ax.scatter(alpha, cv_rmsle_array)
for i, txt in enumerate(np.round(cv_rmsle_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],cv_rmsle_array[i]))

plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha")
plt.ylabel("Error")
plt.show()
```

```

RMSLE for alpha = 1 is 0.45166601980352833
RMSLE for alpha = 2 is 0.44431546233122515
RMSLE for alpha = 3 is 0.4424425182737627
RMSLE for alpha = 3.5 is 0.44171501703551286
RMSLE for alpha = 4 is 0.44154616529145424
RMSLE for alpha = 4.5 is 0.4415286167861061
RMSLE for alpha = 5 is 0.44161632764828285
RMSLE for alpha = 6 is 0.4421832813533032
RMSLE for alpha = 7 is 0.44267468278758176

```



### Training using best hyper-parameters and testing

In [81]:

```

print("Best alpha: ", alpha[best_alpha])
model = Ridge(solver="sag", random_state=42, alpha=alpha[best_alpha])
model.fit(X_train, y_train)
ridge_preds_tr = model.predict(X_train)
ridge_preds_cv = model.predict(X_cv)
ridge_preds_te = model.predict(X_test)

print('Train RMSLE:', sqrt(mse(y_train, ridge_preds_tr)))

ridge_rmsle = sqrt(mse(y_cv, ridge_preds_cv))
print("Cross validation RMSLE: ", ridge_rmsle)

```

```

Best alpha: 4.5
Train RMSLE: 0.3834498984020693
Cross validation RMSLE: 0.4415286167861061

```

## 6.2. Multinomial Naive Bayes: Regression using Classification

[sklearn.naive\\_bayes.MultinomialNB](#)

This model has been used as a correction for the tendency of the Ridge model to underestimate. The predicted values are saved for use as features in a later model. The authors of 18th place solution [kernel](#) improved their RMSLE by 0.003 using this.

In [82]:

```

from sklearn.naive_bayes import MultinomialNB

model = MultinomialNB(alpha=0.01)
model.fit(X_train, y_train>= 4)

mnb_preds_tr = model.predict_proba(X_train)[: , 1]
mnb_preds_cv = model.predict_proba(X_cv)[: , 1]
mnb_preds_te = model.predict_proba(X_test)[: , 1]

```

## SelectKBest: Selecting top 48k features from categorical and text features

In [83]:

```
# from sklearn.feature_selection.univariate_selection import SelectKBest, f_regression

fselect = SelectKBest(f_regression, k=48000)
train_features = fselect.fit_transform(train_sparse, y_train)

cv_features = fselect.transform(cv_sparse)
test_features = fselect.transform(test_sparse)
```

In [84]:

```
print('Shapes after SelectKBest:', train_features.shape, cv_features.shape, test_features.shape)
```

Shapes after SelectKBest: (1332967, 48000) (148108, 48000) (3460725, 48000)

In [85]:

```
# stack feature matrix with Ridge, MNB model predictions, engineered features
X_train = hstack((X_tr.values, ridge_preds_tr.reshape(-1,1), mnb_preds_tr.reshape(-1,1), train_features)).tocsr()

X_cv = hstack((X_val.values, ridge_preds_cv.reshape(-1,1), mnb_preds_cv.reshape(-1,1), cv_features)).tocsr()

X_test = hstack((X_te.values, ridge_preds_te.reshape(-1,1), mnb_preds_te.reshape(-1,1), test_features)).tocsr()

print('Train size: {}, CV size: {}, Test size: {}'.format(X_train.shape, X_cv.shape, X_test.shape))
```

Train size: (1332967, 48049), CV size: (148108, 48049), Test size: (3460725, 48049)

In [86]:

```
del train_features, cv_features
gc.collect()
```

Out[86]:

2685

In [87]:

```
print('Time taken: ', time.time()-start_time)
snapshot2 = tracemalloc.take_snapshot()
top_stats = snapshot2.compare_to(snapshot1, 'lineno')

print("[ Top 10 ]")
for stat in top_stats[:10]:
    print(stat)
```

Time taken: 5293.345500946045

```
[ Top 10 ]
<_array_function__ internals>:6: size=6529 MiB (+6529 MiB), count=520 (+520), average=12.6 MiB
/opt/conda/lib/python3.7/site-packages/pandas/core/algorithms.py:1657: size=1772 MiB (+1772 MiB),
count=13 (+13), average=136 MiB
/opt/conda/lib/python3.7/site-packages/scipy/sparse/compressed.py:755: size=668 MiB (+668 MiB), co
unt=5 (+5), average=134 MiB
<ipython-input-4-1155efdb4829>:4: size=485 MiB (+485 MiB), count=7401015 (+7401015), average=69 B
/opt/conda/lib/python3.7/site-packages/pandas/io/parsers.py:2037: size=400 MiB (+400 MiB),
count=2803386 (+2803386), average=150 B
/opt/conda/lib/python3.7/site-packages/scipy/sparse/compressed.py:754: size=334 MiB (+334 MiB), co
unt=4 (+4), average=83.5 MiB
/opt/conda/lib/python3.7/site-packages/pandas/core/internals/managers.py:1848: size=94.3 MiB (+94.
3 MiB), count=11 (+11), average=8776 KiB
/opt/conda/lib/python3.7/site-packages/sklearn/naive_bayes.py:117: size=75.4 MiB (+75.4 MiB), coun
t=7 (+7), average=10.8 MiB
/opt/conda/lib/python3.7/site-packages/sklearn/naive_bayes.py:117: size=75.4 MiB (+75.4 MiB), coun
t=7 (+7), average=10.8 MiB
```

```
/opt/conda/lib/python3.7/site-packages/scipy/sparse/compressed.py:469: size=38.8 MiB (+38.8 MiB),  
count=10 (+10), average=3977 KiB  
/opt/conda/lib/python3.7/site-packages/pandas/core/algorithms.py:484: size=38.3 MiB (+38.3 MiB), c  
ount=160292 (+160292), average=250 B
```

In [88]:

```
submission_df['price'] = np.exp(ridge_preds_te) - 1  
  
submission_df.to_csv('ridge_submission.csv', index=False)
```

In [89]:

```
scipy.sparse.save_npz("cv_final.npz", X_cv)  
np.save('y_cv', y_cv)  
  
del X_cv, y_cv  
gc.collect()
```

Out[89]:

40

In [90]:

```
scipy.sparse.save_npz("train_final.npz", X_train)  
np.save('y_train', y_train)  
  
del X_train, y_train  
gc.collect()
```

Out[90]:

20

In [91]:

```
scipy.sparse.save_npz("test_final.npz", X_test)  
  
del X_test  
gc.collect()
```

Out[91]:

20